

optimizers-comparison-in-dl

March 2, 2024

```
[1]: #IMPORTING LIBRARIES
```

```
import tensorflow as tf
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D
from tensorflow.keras.layers import Dense, Dropout, Flatten
```

```
2024-03-02 06:26:07.111044: E
external/local_xla/xla/stream_executor/cuda/cuda_dnn.cc:9261] Unable to register
cuDNN factory: Attempting to register factory for plugin cuDNN when one has
already been registered
2024-03-02 06:26:07.111192: E
external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:607] Unable to register
cuFFT factory: Attempting to register factory for plugin cuFFT when one has
already been registered
2024-03-02 06:26:07.257062: E
external/local_xla/xla/stream_executor/cuda/cuda_blas.cc:1515] Unable to
register cuBLAS factory: Attempting to register factory for plugin cuBLAS when
one has already been registered
```

0.0.1 LOADING DATA

```
[2]: (X_train,y_train),(X_test,y_test)=mnist.load_data()
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/mnist.npz
11490434/11490434      0s
0us/step
```

0.0.2 DATA RESHAPE

```
[3]: X_train=X_train.reshape(X_train.shape[0],28,28,1)
X_test=X_test.reshape(X_test.shape[0],28,28,1)
```

```
[4]: X_train=X_train.astype('float32')
X_test=X_test.astype('float32')
X_train /=255
```

```
X_test /=255
y_train=tf.keras.utils.to_categorical(y_train)
y_test=tf.keras.utils.to_categorical(y_test)
```

0.0.3 BUILD OPTIMIZER CALL

```
[5]: def build_optimizer(op):
    model=tf.keras.Sequential()
    model.add(tf.keras.Input(shape=(28,28,1)))
    model.add(tf.keras.layers.Conv2D(filters=32, kernel_size=(3,3), strides=1,
    ↪activation='relu'))
    model.add(tf.keras.layers.MaxPool2D())
    model.add(tf.keras.layers.Conv2D(filters=64, kernel_size=(3,3), strides=1,
    ↪activation='relu'))
    model.add(tf.keras.layers.Dropout(0.25))
    model.add(tf.keras.layers.Flatten())
    model.add(tf.keras.layers.Dense(128, activation='relu'))
    model.add(tf.keras.layers.Dense(256, activation='relu'))
    model.add(tf.keras.layers.Dropout(0.5))
    model.add(tf.keras.layers.Dense(10, activation='softmax'))
    model.compile(optimizer=op, loss='binary_crossentropy',
    ↪metrics=['accuracy'])
    return model
```

0.0.4 COMPARING EACH OPTIMIZER ACCURACY

```
[7]: import os, gc

optimizers=['Adam', 'RMSprop', 'Adadelta', 'Adagrad', 'SGD']
opt_res=[]
model_res=[]
for i in optimizers:
    model=build_optimizer(i)
    print("Accuracy for: ",i)
    print("\n")
    history=model.fit(X_train,y_train, epochs=5, batch_size=64,verbose=1,
    ↪validation_data=(X_test, y_test))
    print("\n")
    gc.collect()
    model_res.append(history)
    opt_res.append(history.history['accuracy'])
```

Accuracy for: Adam

Epoch 1/5
938/938 34s 34ms/step -

accuracy: 0.8167 - loss: 0.0992 - val_accuracy: 0.9815 - val_loss: 0.0103
 Epoch 2/5
 938/938 31s 33ms/step -
 accuracy: 0.9837 - loss: 0.0110 - val_accuracy: 0.9893 - val_loss: 0.0067
 Epoch 3/5
 938/938 31s 33ms/step -
 accuracy: 0.9885 - loss: 0.0076 - val_accuracy: 0.9889 - val_loss: 0.0070
 Epoch 4/5
 938/938 30s 32ms/step -
 accuracy: 0.9920 - loss: 0.0058 - val_accuracy: 0.9900 - val_loss: 0.0066
 Epoch 5/5
 938/938 41s 32ms/step -
 accuracy: 0.9937 - loss: 0.0042 - val_accuracy: 0.9912 - val_loss: 0.0057

Accuracy for: RMSprop

Epoch 1/5
 938/938 31s 32ms/step -
 accuracy: 0.7934 - loss: 0.1095 - val_accuracy: 0.9810 - val_loss: 0.0106
 Epoch 2/5
 938/938 31s 33ms/step -
 accuracy: 0.9780 - loss: 0.0140 - val_accuracy: 0.9865 - val_loss: 0.0077
 Epoch 3/5
 938/938 31s 33ms/step -
 accuracy: 0.9868 - loss: 0.0089 - val_accuracy: 0.9900 - val_loss: 0.0061
 Epoch 4/5
 938/938 42s 34ms/step -
 accuracy: 0.9911 - loss: 0.0061 - val_accuracy: 0.9915 - val_loss: 0.0053
 Epoch 5/5
 938/938 40s 32ms/step -
 accuracy: 0.9939 - loss: 0.0044 - val_accuracy: 0.9900 - val_loss: 0.0063

Accuracy for: Adadelata

Epoch 1/5
 938/938 33s 34ms/step -
 accuracy: 0.1001 - loss: 0.6806 - val_accuracy: 0.1069 - val_loss: 0.6311
 Epoch 2/5
 938/938 31s 33ms/step -
 accuracy: 0.1005 - loss: 0.5991 - val_accuracy: 0.1278 - val_loss: 0.4308
 Epoch 3/5
 938/938 41s 33ms/step -
 accuracy: 0.1077 - loss: 0.4143 - val_accuracy: 0.1536 - val_loss: 0.3345
 Epoch 4/5

938/938 42s 34ms/step -
accuracy: 0.1176 - loss: 0.3604 - val_accuracy: 0.2520 - val_loss: 0.3248
Epoch 5/5
938/938 42s 35ms/step -
accuracy: 0.1237 - loss: 0.3522 - val_accuracy: 0.3556 - val_loss: 0.3196

Accuracy for: Adagrad

Epoch 1/5
938/938 32s 33ms/step -
accuracy: 0.0961 - loss: 0.6229 - val_accuracy: 0.1029 - val_loss: 0.3300
Epoch 2/5
938/938 31s 33ms/step -
accuracy: 0.1181 - loss: 0.3461 - val_accuracy: 0.4434 - val_loss: 0.3171
Epoch 3/5
938/938 30s 31ms/step -
accuracy: 0.1546 - loss: 0.3326 - val_accuracy: 0.6627 - val_loss: 0.3044
Epoch 4/5
938/938 29s 31ms/step -
accuracy: 0.2225 - loss: 0.3183 - val_accuracy: 0.6934 - val_loss: 0.2851
Epoch 5/5
938/938 40s 30ms/step -
accuracy: 0.3294 - loss: 0.2986 - val_accuracy: 0.7023 - val_loss: 0.2536

Accuracy for: SGD

Epoch 1/5
938/938 28s 29ms/step -
accuracy: 0.1162 - loss: 0.5056 - val_accuracy: 0.6243 - val_loss: 0.2982
Epoch 2/5
938/938 42s 30ms/step -
accuracy: 0.3258 - loss: 0.2965 - val_accuracy: 0.7799 - val_loss: 0.1752
Epoch 3/5
938/938 41s 30ms/step -
accuracy: 0.6496 - loss: 0.1944 - val_accuracy: 0.8570 - val_loss: 0.1079
Epoch 4/5
938/938 28s 30ms/step -
accuracy: 0.7635 - loss: 0.1414 - val_accuracy: 0.8869 - val_loss: 0.0815
Epoch 5/5
938/938 30s 32ms/step -
accuracy: 0.8178 - loss: 0.1143 - val_accuracy: 0.9006 - val_loss: 0.0687

0.0.5 PLOTTING OPTIMIZERS ACCURACY

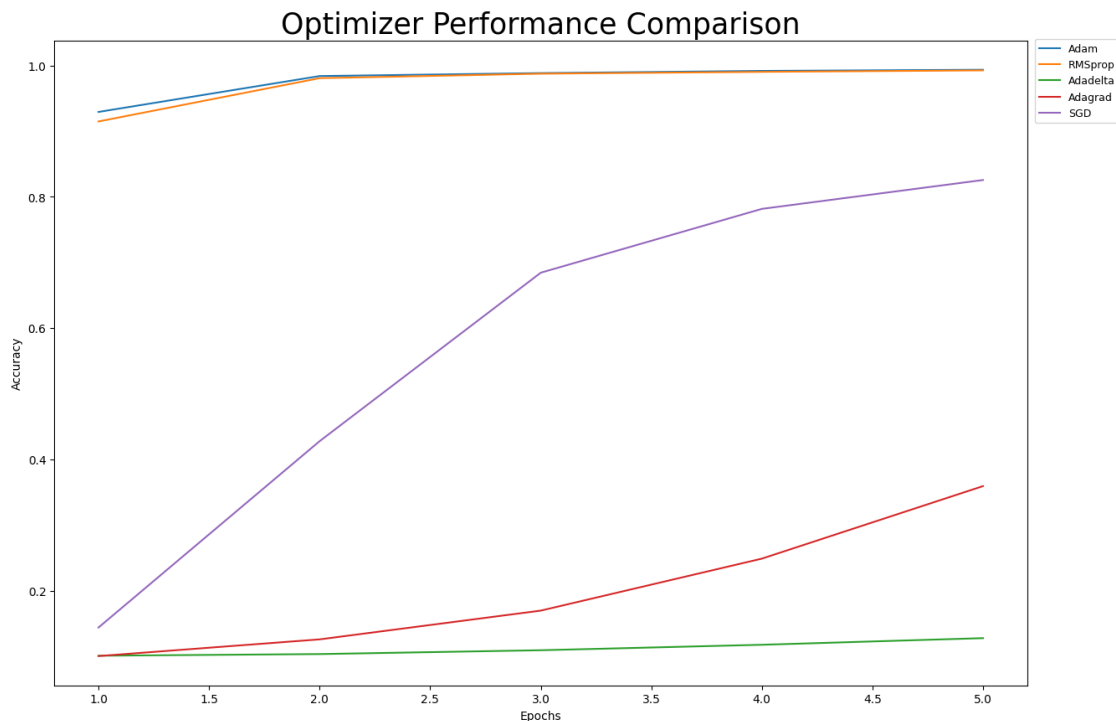
```
[8]: import matplotlib.pyplot as plt

fully_nested = [list(zip(*[(ix+1,y) for ix,y in enumerate(x)])) for x in
    ↪opt_res]
names = ['sublist%d'%(i+1) for i in range(len(fully_nested))]

fig = plt.figure(figsize=(15,10))

for l in fully_nested:
    plt.plot(*l)

plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend(optimizers, fontsize=9, loc = 'upper right', bbox_to_anchor=(1.1, 1.
    ↪01))
plt.title("Optimizer Performance Comparison", fontsize=25)
plt.show()
```



Adam and RMSprop performed the best in terms of accuracy and loss reduction, with Adam being slightly better. Adadelat performed poorly, while Adagrad and SGD showed improvement over epochs but didn't reach the same level of performance as Adam and RMSprop within the given epochs.