

digit-recognizer-cnn

March 6, 2024

```
[65]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
```

```
[66]: train_df = pd.read_csv('/kaggle/input/digitrecognition/train.csv')
test_df = pd.read_csv('/kaggle/input/digitrecognition/test.csv')
```

```
[67]: train_df.head()
```

```
[67]:
```

	label	pixel0	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	\
0	1	0	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	0	0	
2	1	0	0	0	0	0	0	0	0	
3	4	0	0	0	0	0	0	0	0	
4	0	0	0	0	0	0	0	0	0	

	pixel8	...	pixel774	pixel775	pixel776	pixel777	pixel778	pixel779	\
0	0	...	0	0	0	0	0	0	
1	0	...	0	0	0	0	0	0	
2	0	...	0	0	0	0	0	0	
3	0	...	0	0	0	0	0	0	
4	0	...	0	0	0	0	0	0	

	pixel1780	pixel1781	pixel1782	pixel1783
0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0

[5 rows x 785 columns]

```
[68]: train_df.shape
```

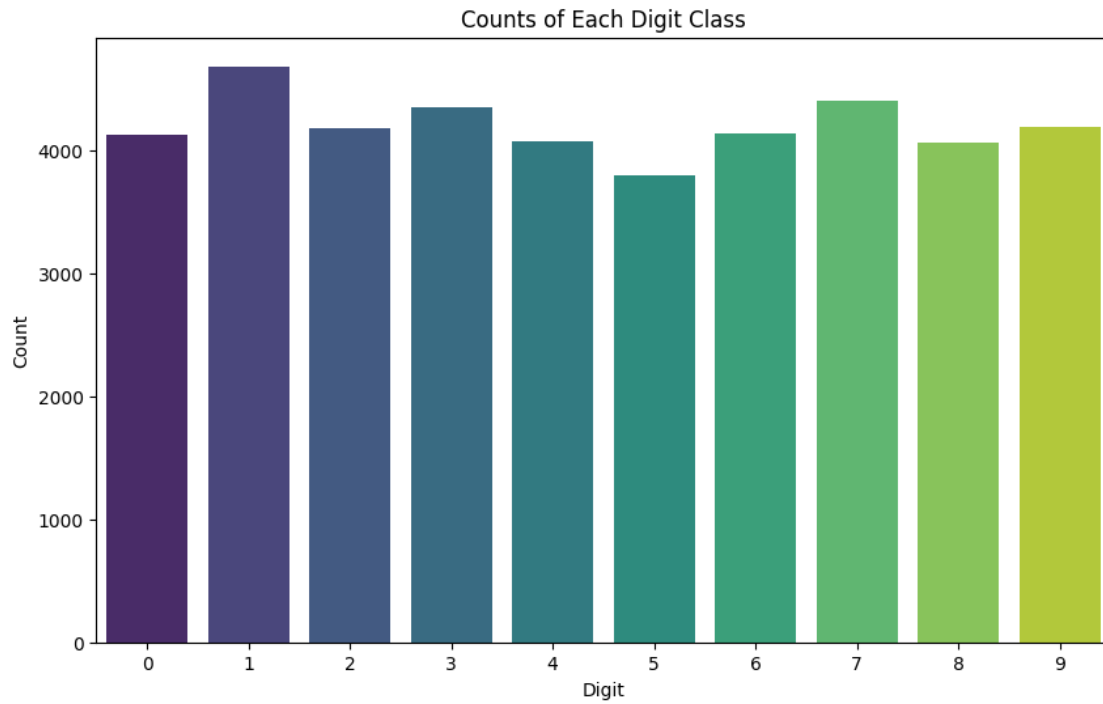
```
[68]: (42000, 785)
```

```
[69]: y = train_df["label"]  
X = train_df.drop(labels=["label"], axis=1)
```

```
[72]: digit_counts = y.value_counts()  
digit_counts
```


```
[72]: label  
1      4684  
7      4401  
3      4351  
9      4188  
2      4177  
6      4137  
0      4132  
4      4072  
8      4063  
5      3795  
Name: count, dtype: int64
```

```
[74]: # Plot the count plot with Viridis palette  
plt.figure(figsize=(10, 6))  
sns.countplot(x=y, palette="viridis")  
plt.title("Counts of Each Digit Class")  
plt.xlabel("Digit")  
plt.ylabel("Count")  
plt.xticks(rotation=0)  
plt.show()
```




```
[76]: # Plot the 2nd sample
img = X.iloc[1].values
img = img.reshape((28,28))
plt.imshow(img, cmap='gray')
plt.title(train_df.iloc[1,0])
plt.axis("off")
plt.show()
```

0



```
[77]: # Plot the 7th sample
img = X.iloc[6].values
img = img.reshape((28,28))
plt.imshow(img, cmap='gray')
plt.title(train_df.iloc[6,0])
plt.axis("off")
plt.show()
```



0.1 Normalization, Reshape and Label Encoding

0.1.1 Normalization

- We perform a grayscale normalization to reduce the effect of illumination's differences.
- If we perform normalization, CNN works faster.

0.1.2 Reshape

- Train and test images (28 x 28)
- We reshape all data to 28x28x1 3D matrices.
- Keras needs an extra dimension in the end which correspond to channels. Our images are gray scaled so it use only one channel.

0.1.3 Label Encoding

- Encode labels to one hot vectors
- 2 => [0,0,1,0,0,0,0,0,0]
- 4 => [0,0,0,0,1,0,0,0,0]

```
[78]: # Normalize the data
X = X / 255.0
test_df = test_df / 255.0
```

```
[79]: # Reshape the data
X = X.values.reshape(-1, 28, 28, 1)
test_df = test_df.values.reshape(-1, 28, 28, 1)
```

```
[80]: from sklearn.preprocessing import LabelEncoder
# Label Encoding
le = LabelEncoder()
y_encoded = le.fit_transform(y)
```

0.2 Train Test Split

We split the data into train and test sets.

test size is 20%.

train size is 80%.

```
[81]: from sklearn.model_selection import train_test_split
X_train, X_val, y_train, y_val = train_test_split(X, y_encoded, test_size=0.2,
↪random_state=2)
```

```
[83]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten, Conv2D, MaxPool2D
from keras.optimizers import Adam
from keras.utils import to_categorical

from keras.utils import to_categorical

# Convert labels to one-hot encoding
Y_train = to_categorical(y_train)
Y_val = to_categorical(y_val)

# Define the model architecture
model = Sequential()
model.add(Conv2D(filters=8, kernel_size=(5, 5), padding='Same',
↪activation='relu', input_shape=(28, 28, 1)))
model.add(MaxPool2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
```

```

model.add(Conv2D(filters=16, kernel_size=(3, 3), padding='Same',
    ↪activation='relu'))
model.add(MaxPool2D(pool_size=(2, 2), strides=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(256, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(10, activation='softmax'))

# Compile the model
model.compile(optimizer=Adam(), loss='categorical_crossentropy',
    ↪metrics=['accuracy'])

# Train the model
history = model.fit(X_train, Y_train, epochs=10, batch_size=64,
    ↪validation_data=(X_val, Y_val), verbose=1)

```

```

Epoch 1/10
525/525          20s 34ms/step -
accuracy: 0.7206 - loss: 0.8540 - val_accuracy: 0.9580 - val_loss: 0.1321
Epoch 2/10
525/525          21s 36ms/step -
accuracy: 0.9381 - loss: 0.1988 - val_accuracy: 0.9708 - val_loss: 0.0909
Epoch 3/10
525/525          20s 35ms/step -
accuracy: 0.9544 - loss: 0.1463 - val_accuracy: 0.9754 - val_loss: 0.0792
Epoch 4/10
525/525          20s 34ms/step -
accuracy: 0.9587 - loss: 0.1266 - val_accuracy: 0.9798 - val_loss: 0.0638
Epoch 5/10
525/525          18s 34ms/step -
accuracy: 0.9648 - loss: 0.1109 - val_accuracy: 0.9801 - val_loss: 0.0611
Epoch 6/10
525/525          18s 34ms/step -
accuracy: 0.9711 - loss: 0.0948 - val_accuracy: 0.9823 - val_loss: 0.0541
Epoch 7/10
525/525          19s 36ms/step -
accuracy: 0.9732 - loss: 0.0850 - val_accuracy: 0.9830 - val_loss: 0.0525
Epoch 8/10
525/525          21s 38ms/step -
accuracy: 0.9744 - loss: 0.0779 - val_accuracy: 0.9856 - val_loss: 0.0445
Epoch 9/10
525/525          19s 36ms/step -
accuracy: 0.9769 - loss: 0.0751 - val_accuracy: 0.9857 - val_loss: 0.0412
Epoch 10/10
525/525          20s 34ms/step -
accuracy: 0.9773 - loss: 0.0712 - val_accuracy: 0.9873 - val_loss: 0.0413


```

0.3 Evaluate the model

Test Loss visualization

Confusion matrix

```
[84]: # Plot the training and validation loss
plt.figure(figsize=(10, 6))
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



```
[86]: # Predict probabilities for each class
y_pred_probs = model.predict(X_val)
y_pred = np.argmax(y_pred_probs, axis=1)
conf_matrix = confusion_matrix(y_val, y_pred)

# Plot the confusion matrix
plt.figure(figsize=(10, 8))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='coolwarm')
plt.title('Confusion Matrix')
```



```
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()
```

263/263

2s 8ms/step

